



JOSÉ RIZAL UNIVERSITY
Find your Inner Hero



Unveiling the Data





DATA PREPROCESSING

Data Preprocessing is an important step of data preparation since datasets often have missing values, incorrect datatypes, and irrelevant features. Through this, the dataset will be cleaned and ready for model training.

- This section consists of:
 - *Preliminary Steps* - In this part, the following are performed: importing essential tools, loading the dataset, constructing the DataFrame, and renaming columns.
 - *Understanding the Dataset* - This part involves performing operations that help us understand the dataset and determine what preprocessing steps must be taken.
 - *Data Cleaning* - Data cleaning was performed by removing duplicates, handling missing values, and ensuring data type consistency and accuracy.



- *Data Transformation* - It involves transforming data through standardization, handling outliers, and extracting components to create new features for improved analysis.
- *Feature Engineering* - Through this, categorical variables were encoded and additional interaction features from existing variables were generated.
- *Data Analysis* - Summary statistics was performed to provide valuable insights about the dataset.
- *Feature Selection and Splitting* - Through correlation, relevant features were selected based on the problem being solved.



Unveiling_the_Data.ipynb ☆ ☁

File Edit View Insert Runtime Tools Help

+ Code + Text

```
[5] 1 import pandas as pd
    2 import numpy as np
    3 import matplotlib.pyplot as plt
```

Essential Libraries:

We'll utilize pandas for data manipulation and analysis, numpy for numerical computations, and matplotlib for data visualization.

```
[70] 1 from google.colab import files
    2 uploaded = files.upload()
```



Choose Files cars.csv

- **cars.csv**(text/csv) - 1816 bytes, last modified: 6/7/2018 - 100% done
- Saving cars.csv to cars (2).csv

from google.colab import files: This line imports the files module specifically designed for Google Colab. This module provides functions for interacting with the file system within the Colab environment, including uploading and downloading files.

uploaded = files.upload(): This is the core of the code. files.upload() is a function within the files module.

Triggers a file selection dialog in your browser: A pop-up window or a section within your Colab notebook will appear, allowing you to browse. Lets you choose one or more files: You select the files you want to upload to your Colab environment. Uploads the selected files: Once you confirm your selection, the files are transferred from your computer to the Colab virtual machine's file



```
✓ [74] 1 df = pd.read_csv("cars.csv")
```

```
✓ [75] 1 df.columns
```

```
Index(['Unnamed: 0', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs',  
      'am', 'gear', 'carb'],  
      dtype='object')
```

Displays the first few rows using `data.head()`.

df.head() in Pandas is a method used to display the first few rows of a DataFrame. It's a quick and easy way to inspect your data and get a sense of its structure and contents.

Explanation:

df: This refers to your Pandas DataFrame.

.head(): This is a method (a function associated with the DataFrame object). Purpose: **df.head()** returns a new DataFrame containing only the first n rows of the original DataFrame. By default, n is 5.

How it works:

1. Retrieves Rows: `df.head()` extracts the specified number of rows from the beginning of the DataFrame.
2. Creates a New DataFrame: It creates a new DataFrame containing only those selected rows. The original DataFrame `df` is not modified.
3. Displays the DataFrame: In an interactive environment like a Jupyter Notebook or Google Colab, `df.head()` often displays the resulting DataFrame directly in a nicely formatted table.



✓
0s



1 `df.head()` # Displays the first 5 rows



	Unnamed: 0	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2



Next steps:

[Generate code with df](#)

[View recommended plots](#)

[New interactive sheet](#)



Specifying the number of rows:

```
1 df.head(10) # Displays the first 10 rows
```



	mpg	hp	gear_category
--	-----	----	---------------



0	21.0	110	4
1	21.0	110	4
2	22.8	93	4
3	21.4	110	3
4	18.7	175	3
5	18.1	105	3
6	14.3	245	3
7	24.4	62	4
8	22.8	95	4
9	19.2	123	4





```
1 df.head(3) # Displays the first 3 rows
```



	Unnamed: 0	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1







df.tail() is very useful for quickly previewing the end of your data. This is often helpful when you've just added new data to your DataFrame or when you want to check the last few entries. Combined with **df.head()**, you can get a good overview of your data's structure and content without having to display the entire DataFrame, especially when it's large.

 1 df.tail()



	Unnamed: 0	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.9	1	1	5	2
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.5	0	1	5	4
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.5	0	1	5	6
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.6	0	1	5	8
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.6	1	1	4	2





JOSÉ RIZAL UNIVERSITY
Find your Inner Hero



data.shape in Pandas (and also in NumPy) is an attribute that returns a tuple representing the dimensions of your DataFrame (or array). It tells you the number of rows and columns in your data.

Return Value: data.shape returns a tuple of two integers: (number_of_rows, number_of_columns).

Get an overview of data dimensions (rows and columns) with data.shape.

```
[79] 1 df.shape
```

```
↔ (32, 12)
```





df.isnull().sum() is a powerful combination of Pandas methods used to count the number of missing (or null) values in each column of a DataFrame.

isnull(): This method, when applied to a DataFrame, returns a new DataFrame of the same shape as the original. Each cell in the new DataFrame contains a boolean value:

True if the corresponding cell in the original DataFrame contains a missing value (NaN, None, etc.).

False if the cell contains a non-missing value.

✓
0s

1 df.isnull().sum()



0

mpg

0

hp

0

gear_category

0

dtype: int64



df.info() in Pandas is a method that provides a concise summary of information about your DataFrame. It's a very useful way to get a quick overview of your data's structure and characteristics.

.info(): This is a method (a function associated with the DataFrame object). Purpose: df.info() prints information about the DataFrame to the console (or output area).

What information does df.info() provide?

DataFrame Class: It tells you that you're working with a Pandas DataFrame.

RangeIndex: It shows the index range (row labels). Often, this is a simple numerical range like RangeIndex: 8 entries, 0 to 7 if you have 8 rows.

Data columns (total n columns): It lists all the columns in your DataFrame, along with:

Column Name: The name of each column.

Non-Null Count: The number of non-missing (non-null) values in each column. This is very helpful for identifying missing data.

Dtype: The data type of each column (e.g., int64, float64, object, bool). object often indicates string data or mixed data types.

Memory Usage: It shows the amount of memory the DataFrame is using. This can be useful for understanding memory efficiency, especially w



```
1 df.info()
```



```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 32 entries, 0 to 31  
Data columns (total 12 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   Unnamed: 0   32 non-null     object  
1   mpg          32 non-null     float64  
2   cyl          32 non-null     int64  
3   disp         32 non-null     float64  
4   hp           32 non-null     int64  
5   drat         32 non-null     float64  
6   wt           32 non-null     float64  
7   qsec         32 non-null     float64  
8   vs           32 non-null     int64  
9   am           32 non-null     int64  
10  gear         32 non-null     int64  
11  carb         32 non-null     int64  
dtypes: float64(5), int64(6), object(1)  
memory usage: 3.1+ KB
```





DATA TYPES

`df.dtypes` in Pandas is an attribute that returns a Series containing the data type of each column in the DataFrame `df`. It's a quick way to see the data types of all your columns at once.

dtypes: This is an attribute (not a method/function). You access it directly without parentheses.

Return Value: `df.dtypes` returns a Pandas Series. The index of the Series is the column names of the DataFrame, and the values of the Series are the corresponding data types of those columns.

1 `df.dtypes`



0

Unnamed: 0 object

mpg float64

cyl int64

disp float64

hp int64

drat float64

wt float64

qsec float64

vs int64





Descriptive Statistics:

Summarize Statistics



```
1 data.describe()
```



	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
count	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.0000
mean	20.090625	6.187500	230.721875	146.687500	3.596563	3.217250	17.848750	0.437500	0.406250	3.687500	2.8125
std	6.026948	1.785922	123.938694	68.562868	0.534679	0.978457	1.786943	0.504016	0.498991	0.737804	1.6152
min	10.400000	4.000000	71.100000	52.000000	2.760000	1.513000	14.500000	0.000000	0.000000	3.000000	1.0000
25%	15.425000	4.000000	120.825000	96.500000	3.080000	2.581250	16.892500	0.000000	0.000000	3.000000	2.0000
50%	19.200000	6.000000	196.300000	123.000000	3.695000	3.325000	17.710000	0.000000	0.000000	4.000000	2.0000
75%	22.800000	8.000000	326.000000	180.000000	3.920000	3.610000	18.900000	1.000000	1.000000	4.000000	4.0000
max	33.900000	8.000000	472.000000	335.000000	4.930000	5.424000	22.900000	1.000000	1.000000	5.000000	8.0000





Working with .columns

Getting Column Names: The most common use is to simply get the names of the columns:

```
1 column_names = df.columns
2 print(column_names) # Output will be the Index object
3 print(list(column_names)) # To get a regular Python list of column names
4
5 for col in df.columns:
6     print(col) # Print each column name individually
```

```
Index(['Unnamed: 0', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs',
      'am', 'gear', 'carb'],
      dtype='object')
['Unnamed: 0', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am', 'gear', 'carb']
Unnamed: 0
mpg
cyl
disp
hp
drat
wt
qsec
vs
am
gear
carb
```



ISO 21001:2018

TUVNORD



Selecting Columns: You can use `df.columns` along with other Pandas methods to select specific columns:

```
js 1 columns_to_keep = ['mpg', 'hp', 'gear']  
2 df_subset = df[columns_to_keep] # or df = df[[col for col in df.columns if col in columns_to_keep]]
```

**Displays the first few rows **

```
js [30] 1 df_subset.head()
```




	mpg	hp	gear
0	21.0	110	4
1	21.0	110	4
2	22.8	93	4
3	21.4	110	3
4	18.7	175	3





Renaming Columns: While you don't directly modify `df.columns`, you can use it in conjunction with the `df.rename()` function to rename columns:

```
1 df = df_subset.rename(columns={'gear': 'gear_category'})
2 df
```



	mpg	hp	gear_category
0	21.0	110	4
1	21.0	110	4
2	22.8	93	4
3	21.4	110	3
4	18.7	175	3
5	18.1	105	3
6	14.3	245	3

